

Problem 1.

1. acc: 44.15 +- 0.83%

epochs: 80, **each epoch:** 300 batches (80*300 episodes)

distance function: euclidean, **optimizer:** Adam, **lr:** 1e-4

data augmentation: RandomResizedCrop,
RandomHorizontalFlip, RandomGrayScale

meta-train: 25-way 1-shot, **meta-val:** 5-way 1-shot

query: 15 samples / batch

```
Convnet(
  (encoder): Sequential(
    (0): Sequential(
      (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (2): ReLU()
      (3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    )
    (1): Sequential(
      (0): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (2): ReLU()
      (3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    )
    (2): Sequential(
      (0): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (2): ReLU()
      (3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    )
    (3): Sequential(
      (0): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (2): ReLU()
      (3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    )
  )
  (MLP): MLP(
    (mlp): Sequential(
      (0): Linear(in_features=1600, out_features=400, bias=True)
      (1): Dropout(p=0.5, inplace=False)
      (2): ReLU()
      (3): Linear(in_features=400, out_features=64, bias=True)
    )
  )
)
```

2. for parametric distance method:

I **expand** the **outputs** of the support set and query set to the same shape, and **concatenate** them together.

It'll get the shape (**support num, query num, 64+64**)

Then do **MLP** to get the similarity of the instances.

```
(distance_nn): MLP(
  (mlp): Sequential(
    (0): Linear(in_features=128, out_features=32, bias=True)
    (1): Dropout(p=0.5, inplace=False)
    (2): ReLU()
    (3): Linear(in_features=32, out_features=1, bias=True)
  )
)
```

meta-train: 5-way 1-shot

meta-test: 5-way 1-shot

| distance method | valid acc. |
|-------------------|----------------|
| euclidean | 43.16 +- 0.82% |
| cosine | 19.88 +- 0.31% |
| parametric | 39.51 +- 0.73% |

I didn't get a good performance in the parametric method, but it was good at my sample method (random sample 300 batches).

One possibility is that my sample method is not good enough.

The other possibility may be the

3. lr: 1e-4 / epochs: 80 / episodes: 80*300

| meta-train/test | distance method | valid acc. |
|----------------------|-----------------|------------|
| 5-way 1-shot | euclidean | 43.22 |
| 5-way 5-shot | euclidean | 63.27 |
| 5-way 10-shot | euclidean | 67.19 |

In my problem 3, the valid samples are based on my sampler.

My 5-way 5/10-shot seems to have a high acc. at the valid set, however, they're not perform well at meta-test in 5-way 1-shot. Can't even reach 40% accuracy.

It seems reasonable to get a higher acc. with more shots in both training and testing with more samples for the support center.

Problem 2.

1. When pretraining the SSL model, I used the default BYOL method to pretrain the resnet50(img size=128).

I only used the RandomHorizontalFlip(p=20) in augmentation, others are wrapped in BYOL .

Optimizer: Adam(lr=3e-4)

Batch: 64

Batches: 40 (about 24hr)

Of course, without labels.

2.

| setting | Pre-training | Fine-tuning | valid acc. |
|---------|--------------|-----------------|------------|
| A | - | full model | 37.19 |
| B | TA's | full model | 47.29 |
| C | My SSL | full model | 40.15 |
| D | TA's | classifier only | 24.88 |
| E | MySSL | classifier only | 18.97 |

3. In my problem 2, the setting B gets the highest accuracy of the validation set, however, the loss is high. So it may be a little bit of

overfitting in the office dataset.

The setting D also got a high loss.

The setting A and C have similar loss, but with the pretrained model, the setting C had higher accuracy(3%). So the pretrained model seems helpful.